

FIFO Queues

are **all** You Need for Cache Eviction

Juncheng Yang

Yazhuo Zhang, Ziyue Qiu, Yao Yue, K. V. Rashmi

**Carnegie
Mellon
University**

Carnegie Mellon
Parallel Data Laboratory



Software cache and eviction

- Ubiquitous deployments of software caches
 - page cache, block cache, database cache
 - key-value cache, object cache...
- Cache metrics
 - efficiency / effectiveness: miss ratio
 - throughput and scalability: requests/sec
 - simplicity
- A core component of cache design: **eviction**



Cachelib



Pelikan



A long history of research centered around LRU

- Least-recently-used (LRU)
 - maintains objects in a queue with last-access order
 - metadata update (with locking) on each read request
- Problems
 - not scalable
 - not scan-resistant



A long history of research centered around LRU

- Improve LRU's efficiency
 - **add more techniques/queues/metrics:** ARC^[FAST'03], LIRS^[SIGMETRICS'02], LRU-K^[SIGMOD'93], 2Q^[VLDB'94], MQ^[ATC'01], TinyLFU^[TOS'17], LRB^[NSDI'20], CACHEUS^[FAST'21]...
- Improve LRU's throughput and scalability
 - **reduce #operations per-request:** relaxed LRU, CLOCK variants
- **Existing works: tradeoff between efficiency and throughput**

An alternative: FIFO eviction algorithm

- First-in-first-out (FIFO)
 - simpler
 - fewer metadata
 - less computation
 - more scalable
 - flash friendly



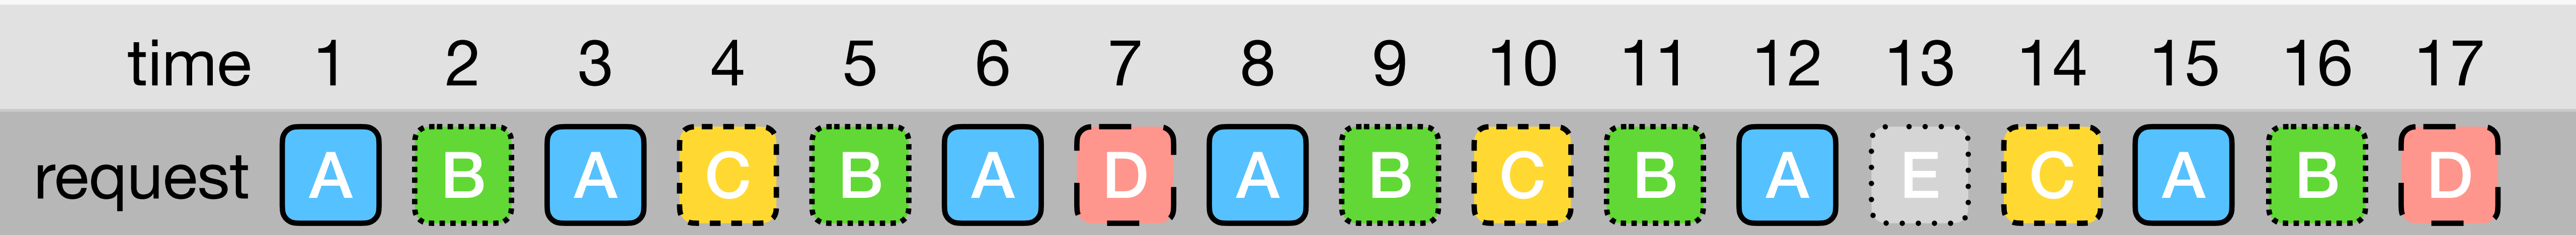
The only drawback:
FIFO has a high miss ratio

**Can we design an efficient
FIFO-based algorithm?**

Observation

More one-hit wonders than you would have expected

- One-hit wonder: objects appeared once in the sequence
- Zipfian workloads: One-hit-wonder ratio **decreases with sequence length**
(measured in #obj)

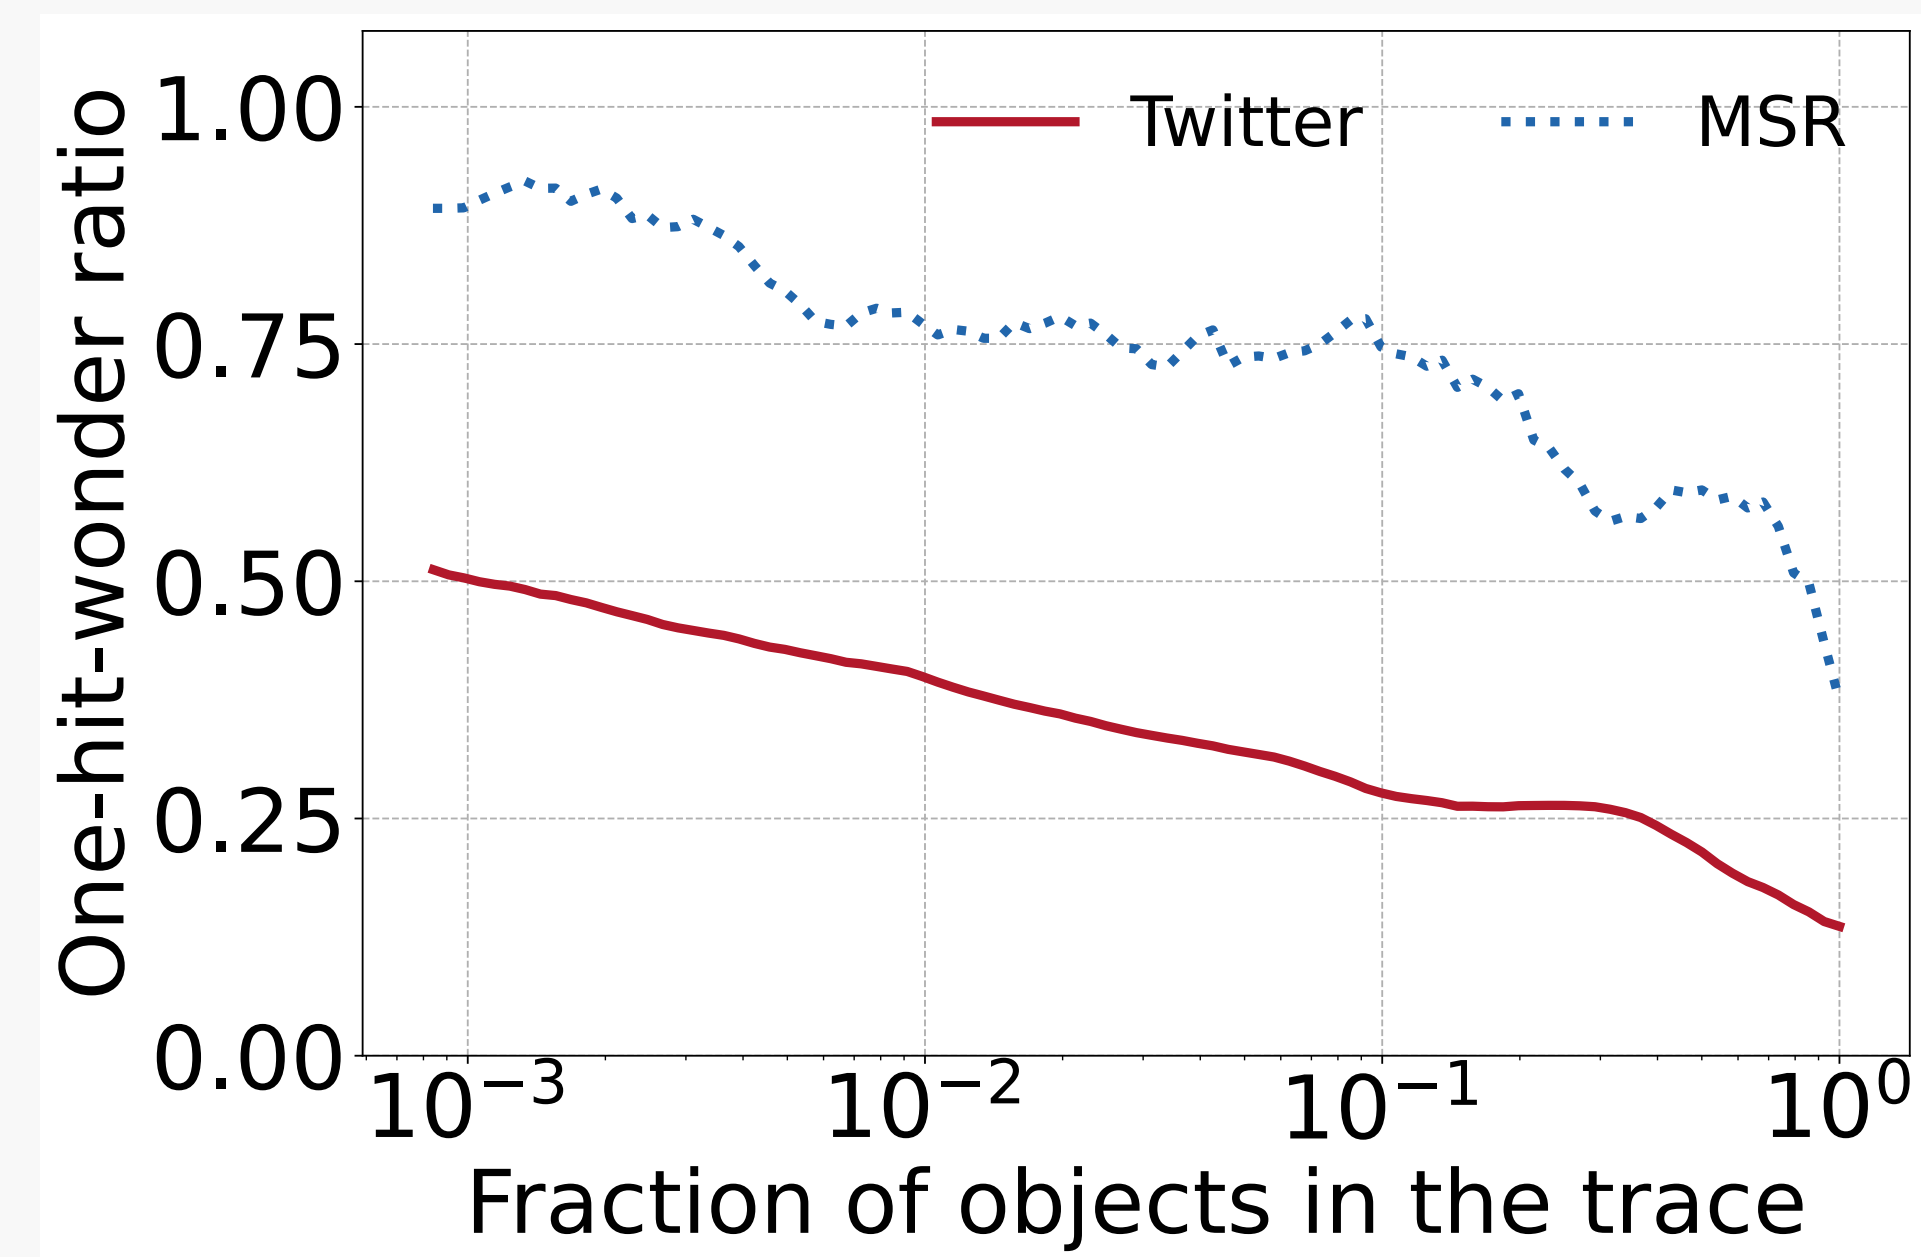
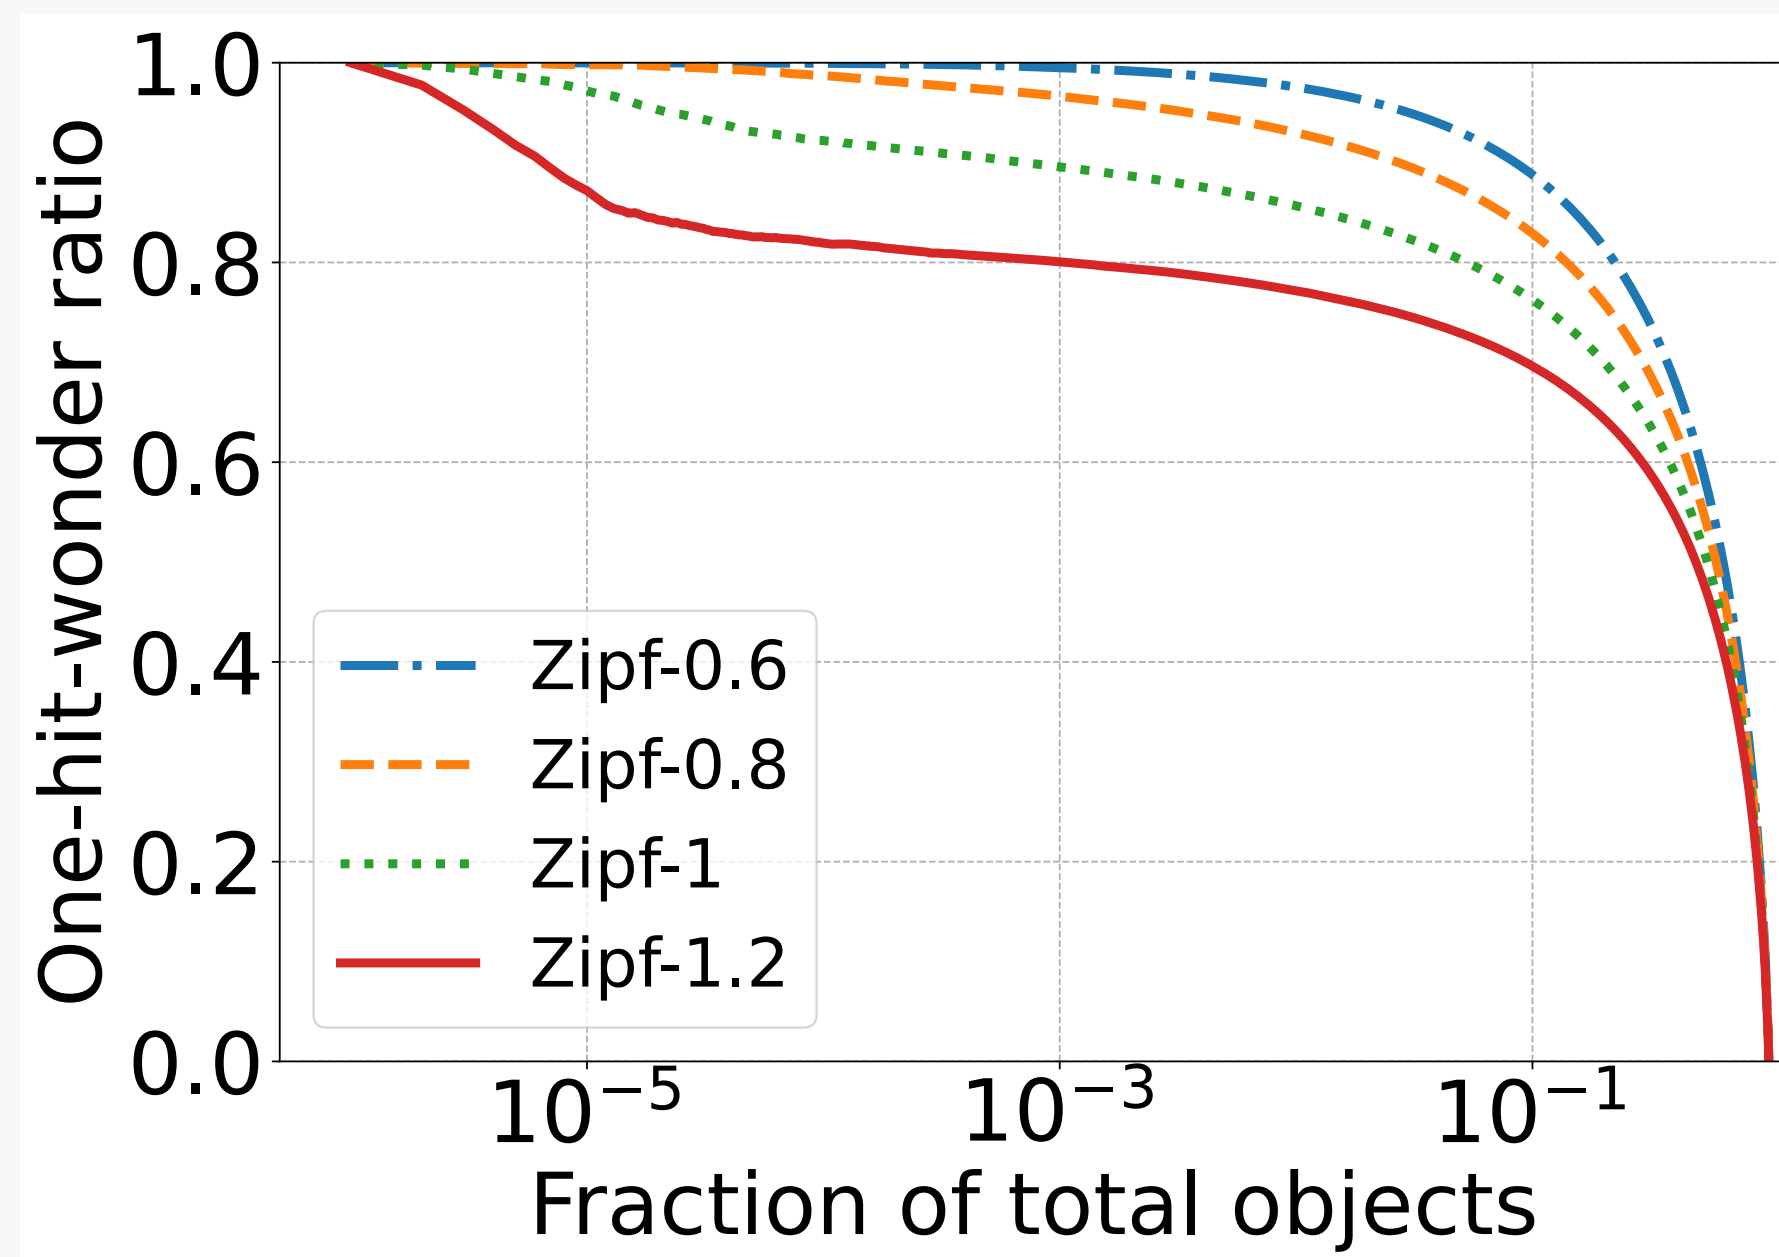


start time	end time	sequence length (# objects)	# one-hit wonder	one-hit wonder ratio
1	17	5	1 (E)	20%
1	7	4	2 (C, D)	50%
1	4	3	2 (B, C)	66%

Observation


More one-hit wonders than you would have expected

- One-hit wonder: objects appeared once in the sequence
- Zipfian workloads: One-hit-wonder ratio **decreases with sequence length** (measured in #obj)



Observation


More one-hit wonders than you would have expected



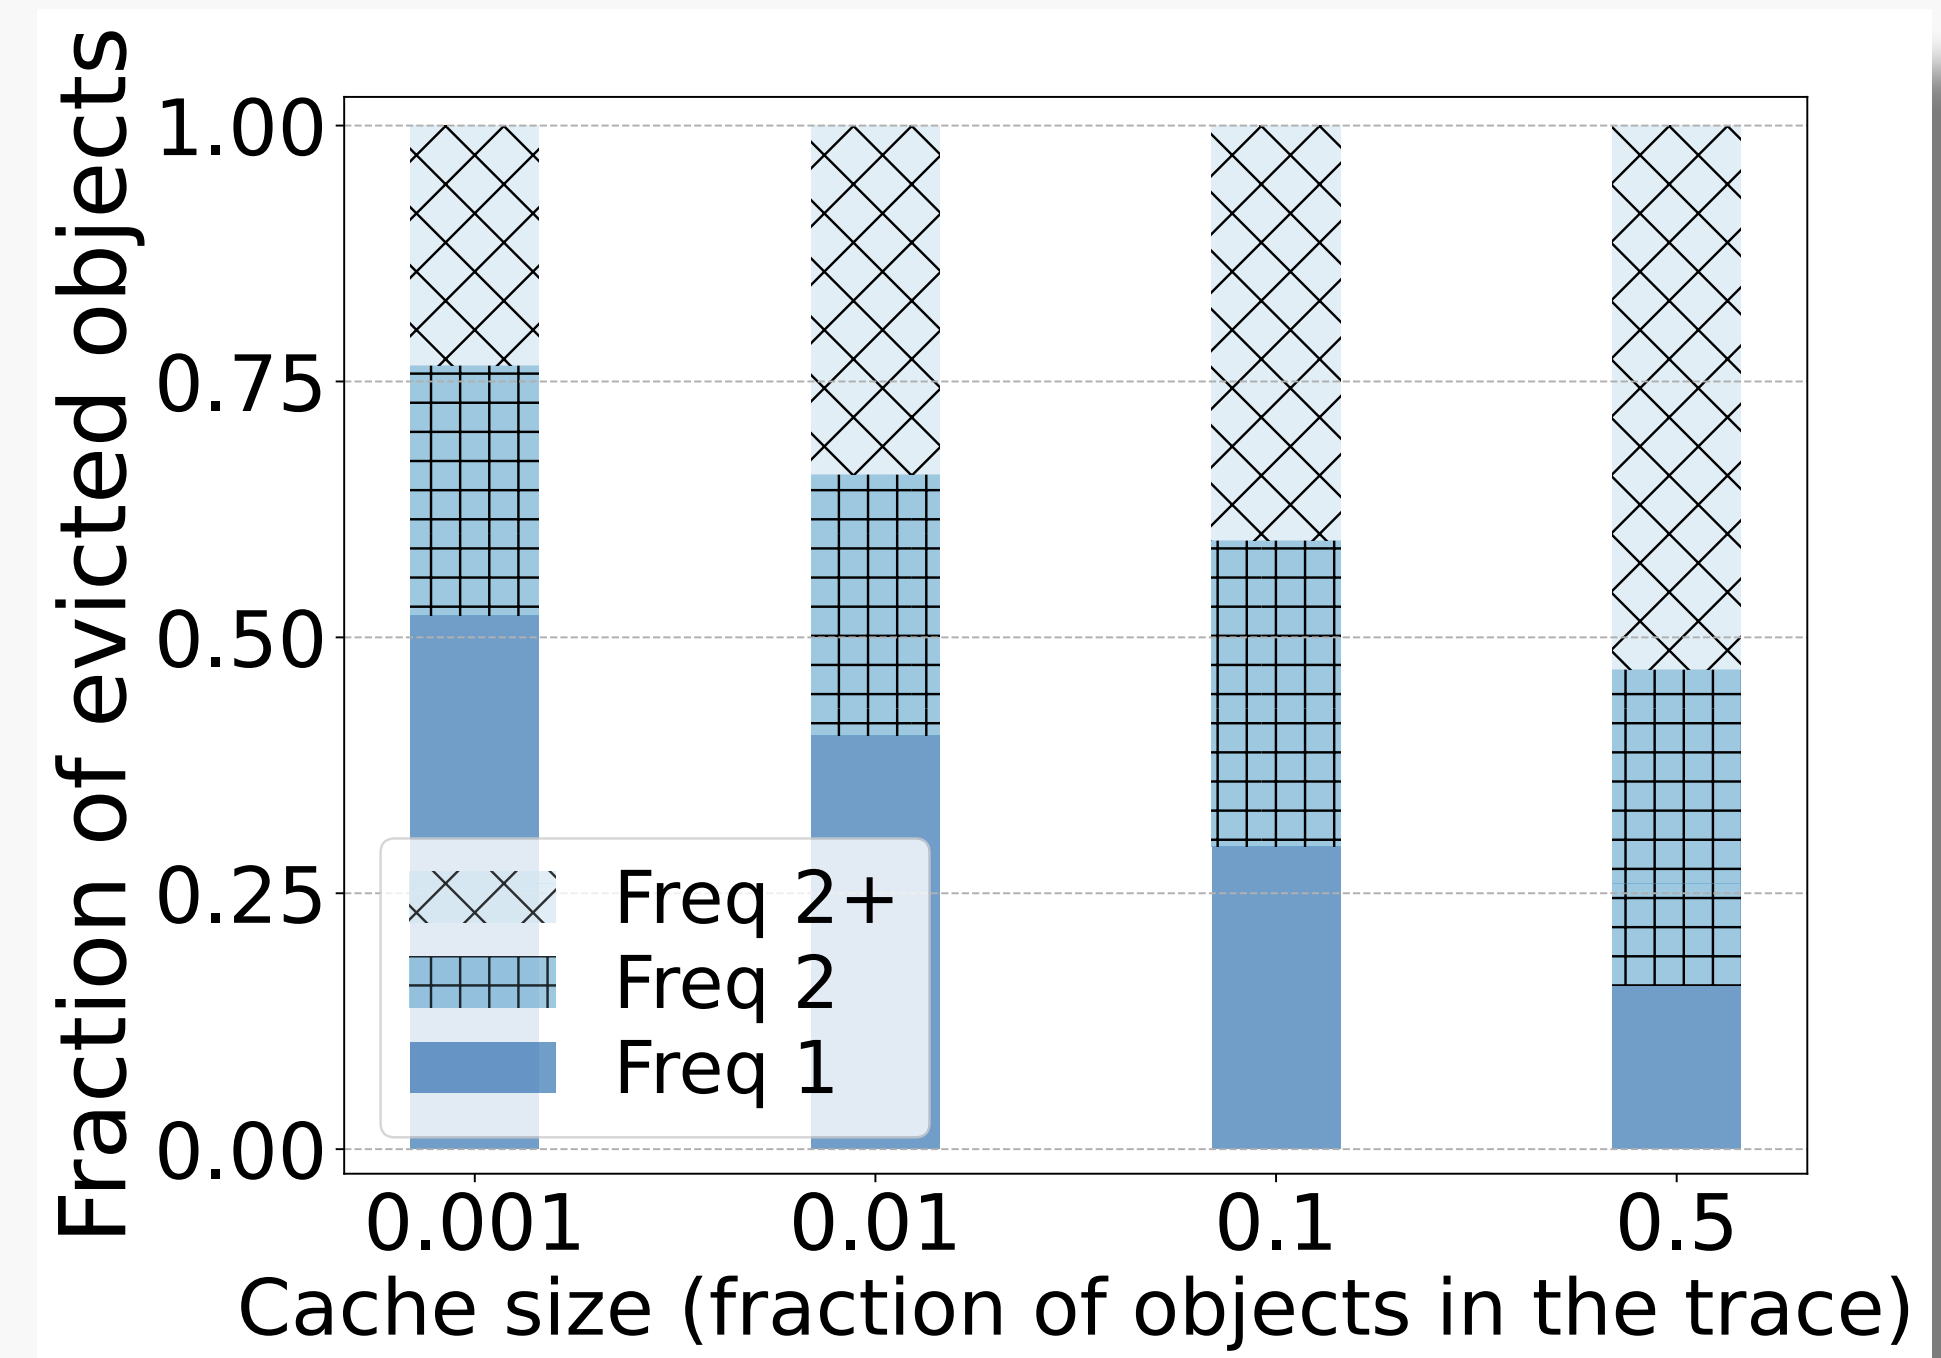
A cache starts eviction after a short request sequence
=> **most objects in the cache are one-hit wonders**

Observation

Many objects in the cache are one-hit-wonders



LRU cache running MSR workload



LRU cache running Twitter workload

A cache starts eviction after a short request sequence
=> **most objects in the cache are one-hit wonders**

S3-FIFO Design

Simple, Scalable eviction algorithm with three Staic FIFO queues



<https://s3fifo.com>

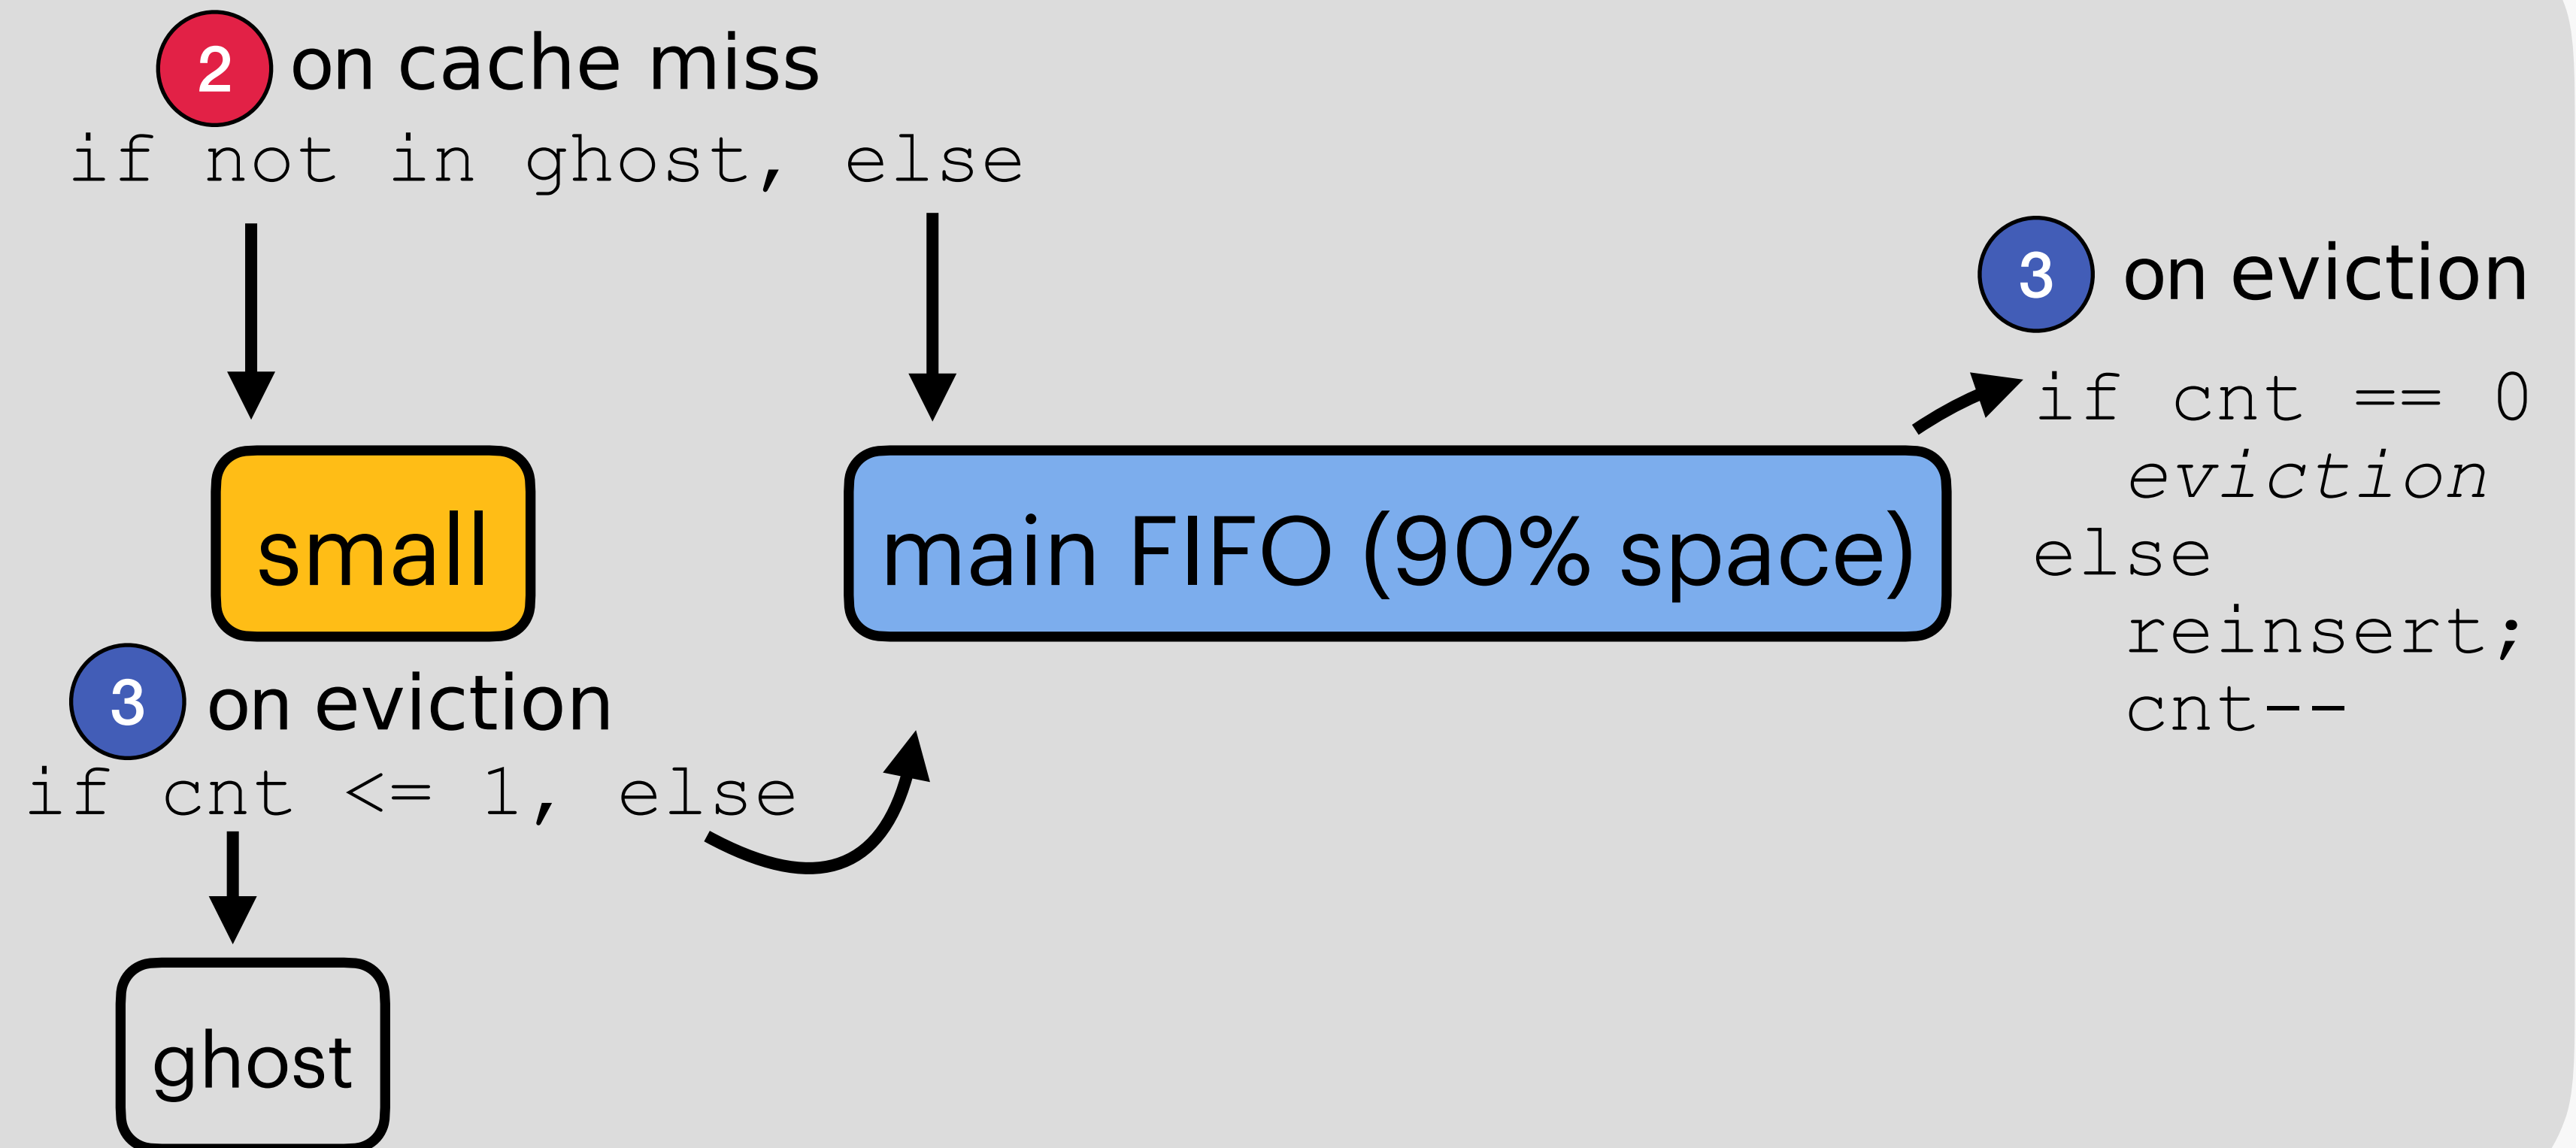
<https://blog.s3fifo.com>

S3-FIFO design

Simple, Scalable eviction algorithm with three Static FIFO queues

```
struct object {  
    ...  
    uint8_t cnt:2;  
}
```

1 on cache hit
cnt++

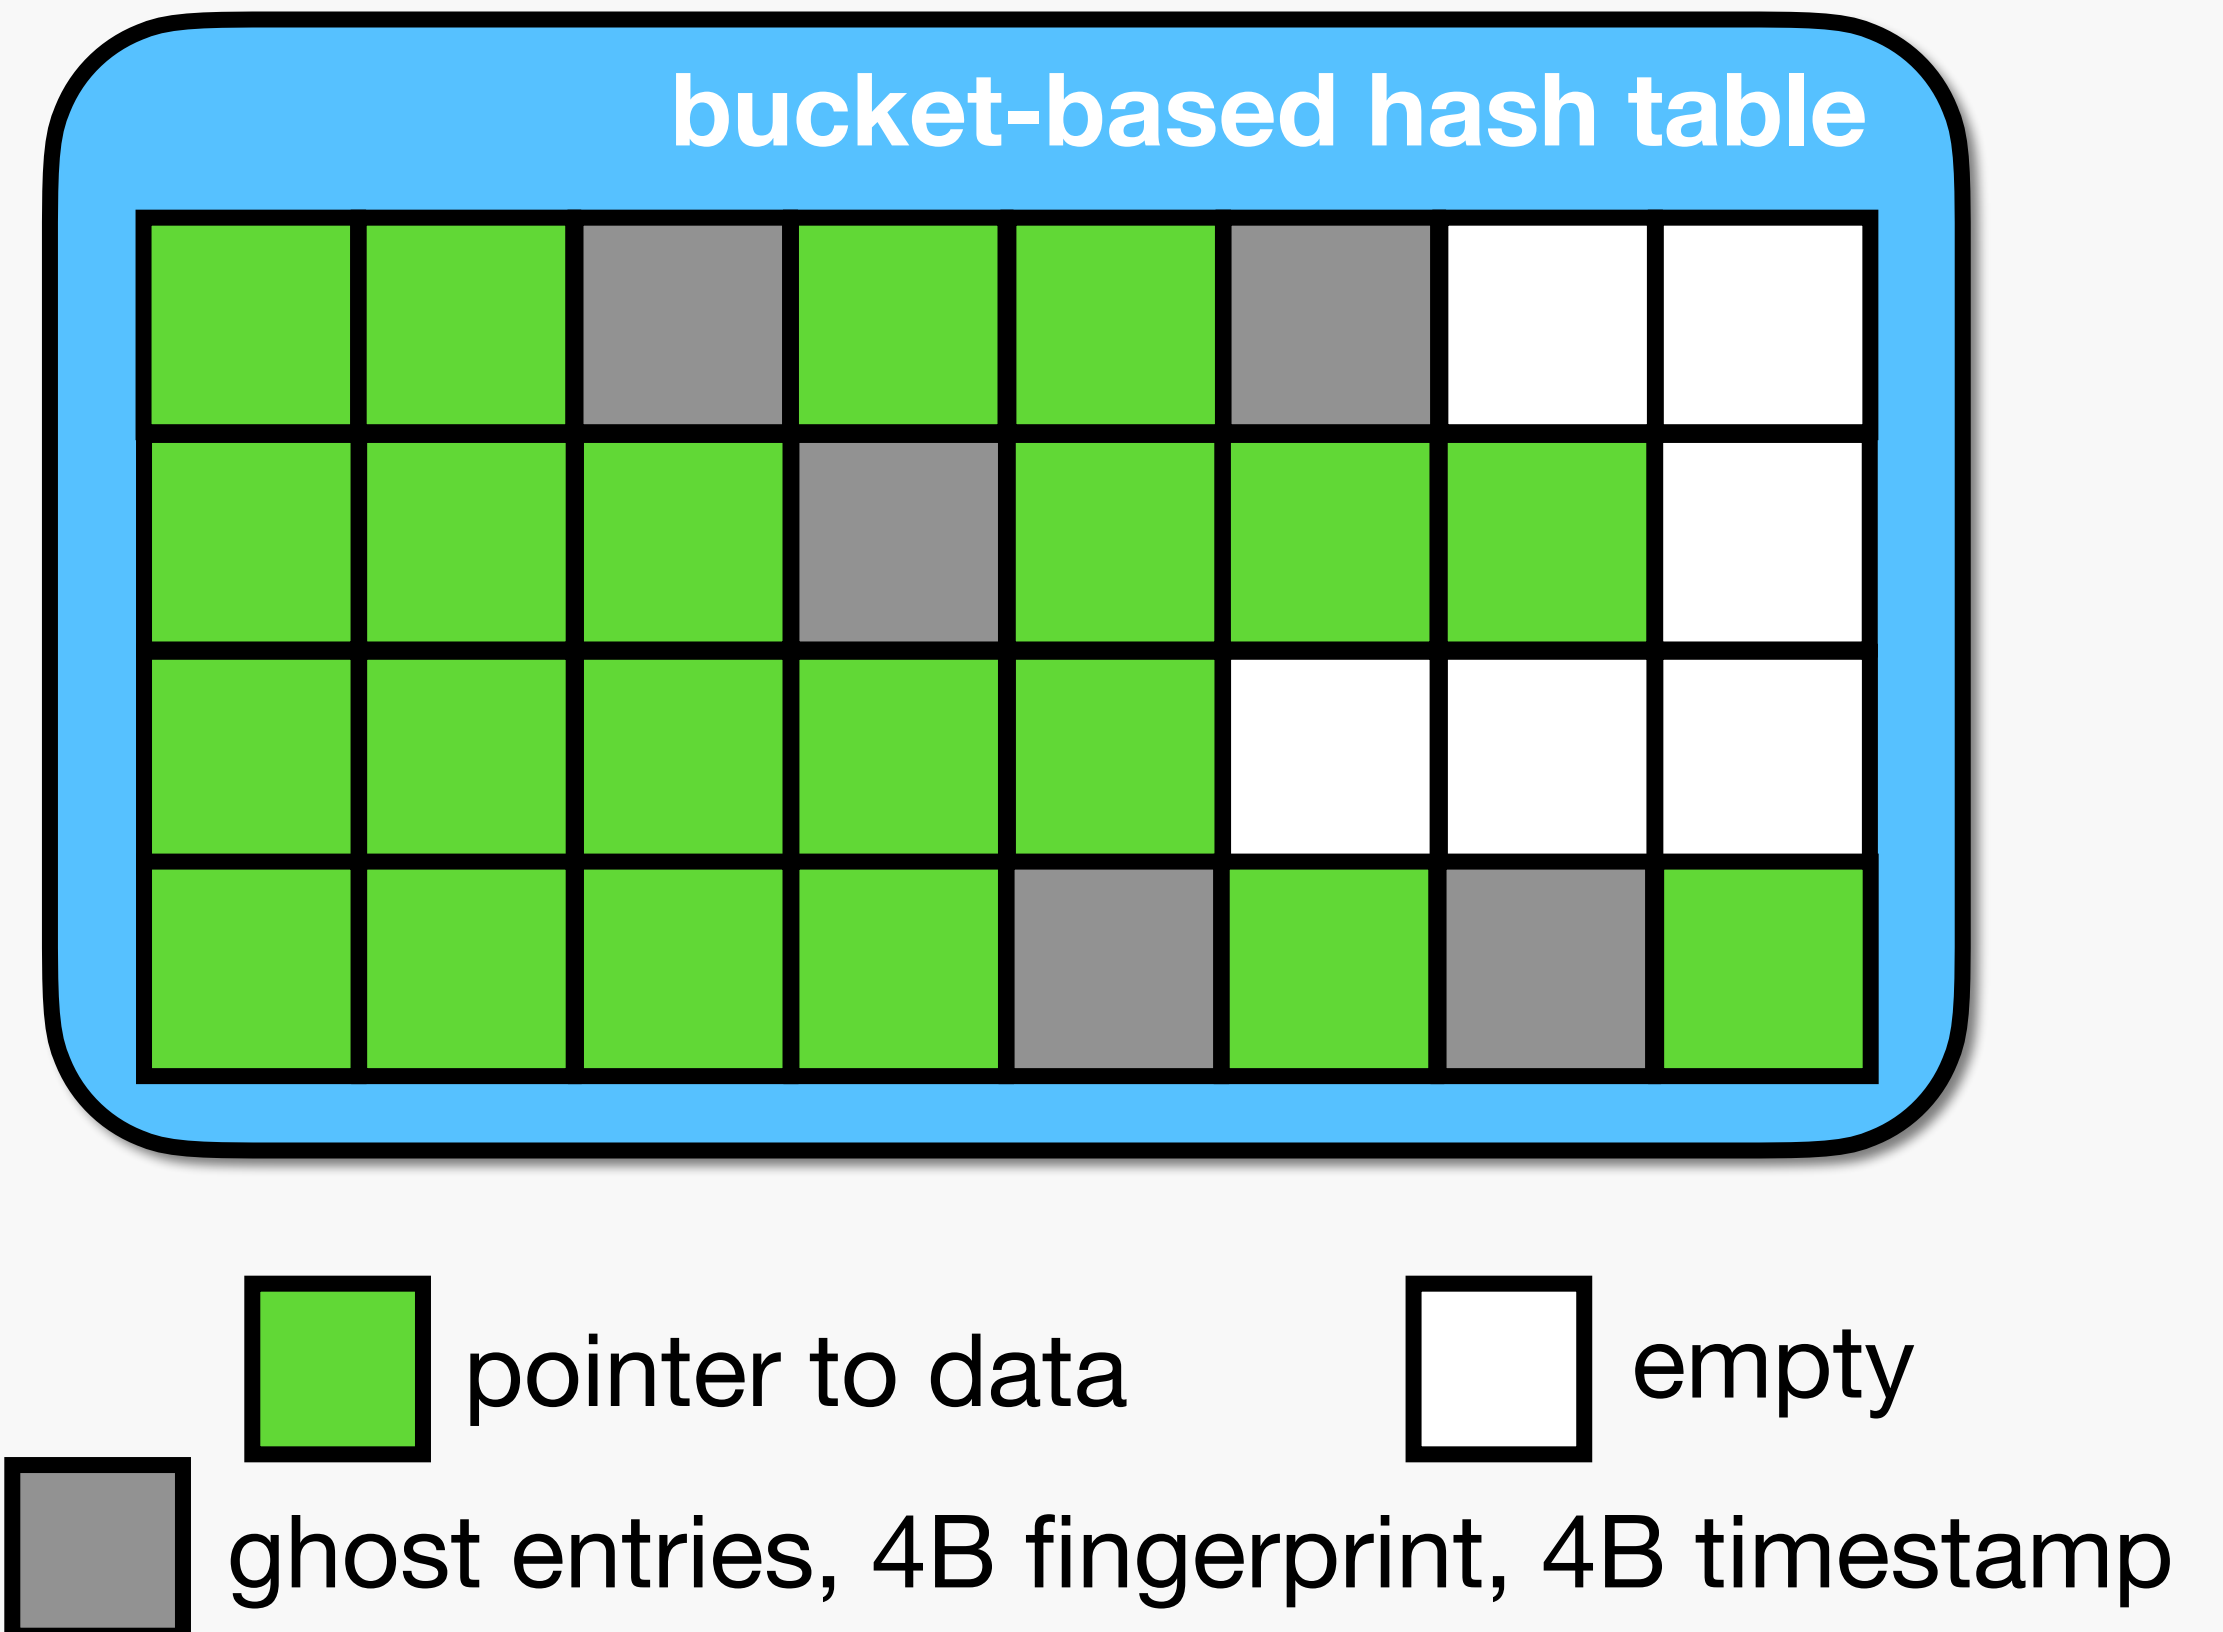


S3-FIFO features

- **Simple and robust:** static queues
- **Fast:** no metadata update for most requests
- **Scalable:** no lock
- **Tiny metadata:** 2 bits
- **Flash friendly:** sequential writes

S3-FIFO implementation

- Can implement with one, two or three FIFO queues
- The small and main FIFO queues can be merged
- Ghost FIFO can be implemented as part of the index (e.g., hash table)
 - in the queue $\Rightarrow T_{curr} - T_{insert} < S_{ghost}$



S3-FIFO evaluation

Evaluation setup


- Dataset
 - 14 datasets, 6594 traces from Twitter, Meta, Microsoft, Wikimedia, Tencent, Alibaba, major CDNs...
 - 848 billion requests , 60 billion objects
 - collected between 2007 and 2023
 - block, key-value, object caches
- Platform
 - libCacheSim, cachelib
 - in-house distributed computation system
- Metric
 - miss ratio reduction from FIFO
 - throughput



CloudLab

Efficiency


Miss ratio reduction distribution across all traces (large size)



Better than state-of-the-art algorithms, up to 72% lower miss ratio than LRU

Efficiency



Mean miss ratio reduction on each dataset (large size)



- **efficient**: the largest mean miss ratio reduction or is close to the best
- **robust**: the best on 10 of the 14 datasets

Throughput and scalability

Zipf workloads




- the fastest on a single thread
- more scalable than optimized LRU, 6x higher throughput
- close to Segcache^[NSDI'21]


More in the paper

- Why S3-FIFO is effective
- Implication for flash cache
- Byte miss ratio
- Impact of small FIFO size
- What if we use LRU

Takeaway




- **Cache workloads exhibit high one-hit-wonder ratio**
 - many objects in the cache are not re-accessed
 - critical to remove the one-hit wonders early
- **S3-FIFO**: simple, scalable caching with **three** static **FIFO** queues
 - reinsertion to keep popular objects
 - a small FIFO queue to *quickly* filter out one-hit wonders
 - interests from Google, VMWare vSAN, Cloudflare, TigerBeetle, Kuaishou, etc.



[HotOS'23] FIFO queues
can be better than LRU

juncheny@cs.cmu.edu



<https://s3fifo.com>
<https://blog.s3fifo.com>

